

## OpenMP Pragma Cheat-Sheet

Run the code: `g++ -fopenmp prog.cpp -o prog`

`#include <omp.h> // This is the header`

| Pragma                                | Purpose   |
|---------------------------------------|---|
| <code>#pragma omp parallel</code>     | Starts a parallel region where multiple threads execute the enclosed code concurrently. |
| <code>#pragma omp for</code>          | Distributes loop iterations among threads within a parallel region.                     |
| <code>#pragma omp parallel for</code> | Combines parallel and for to create a parallel loop directly.                           |
| <code>#pragma omp single</code>       | Specifies a block of code to be executed by only one thread.                            |
| <code>#pragma omp critical</code>     | Defines a critical section where only one thread can execute the block at a time.       |
| <code>#pragma omp barrier</code>      | Synchronizes all threads, making each wait until all have reached the barrier.          |
| <code>#pragma omp master</code>       | Specifies a block of code to be executed only by the master thread.                     |
| <code>#pragma omp sections</code>     | Divides work into separate sections to be executed by different threads.                |
| <code>#pragma omp section</code>      | Defines a single section within a sections block.                                       |
| <code>#pragma omp task</code>         | Creates a task that can be executed by any thread in the team.                          |
| <code>#pragma omp atomic</code>       | Ensures atomic updates to a variable, preventing race conditions with minimal overhead. |
| <code>#pragma omp reduction</code>    | Performs a reduction operation (e.g., sum, product) across all threads.                 |

## Common Clauses in OpenMP Pragmas

OpenMP pragmas can include various clauses to control the behavior of the parallel regions and loops. Some common clauses include:

- **private(variable):** Each thread has its own instance of the variable.
- **shared(variable):** The variable is shared among all threads.
- **default(none):** Requires explicit specification of the data-sharing attributes for all variables.
- **firstprivate(variable):** Each thread has its own instance of the variable, initialized with the value before the parallel region.
- **lastprivate(variable):** The value of the variable from the last iteration is copied back to the original variable after the parallel region.
- **schedule(kind, chunk\_size):** Controls how loop iterations are divided among threads. Common kind values include static, dynamic, guided, and auto.

You can schedule in two ways, one is clause-way:

```
#pragma omp parallel for schedule(static, chunk_size)
#pragma omp parallel for schedule(dynamic, chunk_size)
#pragma omp parallel for schedule(guided, chunk_size)
```

Other way is like mentioning as function:

```
omp_sched_t sched_type // is the parameter
(takes values omp_sched_static/omp_sched_dynamic/..)
```

```
omp_set_schedule(sched_type, chunk_size)
```

- **static:** Iterations are divided into chunks of chunk\_size and assigned to threads in a round-robin fashion. (omp\_sched\_static)
- **dynamic:** Threads request new chunks as they finish their current ones. (omp\_sched\_dynamic)
- **guided:** Similar to dynamic but with decreasing chunk sizes. (omp\_sched\_guided)
- **auto:** The scheduling decision is left to the compiler/runtime. (omp\_sched\_auto)
- **collapse(n):** Merges n nested loops into a single loop for parallelization.
- **reduction(operator:variable):** Specifies a reduction operation on a variable.

## How to use Pragma for recursion: (Sample program: nth fibonacci)

```
#include <omp.h>
#include <iostream>
```

```

int fib(int n) {
    int left, right;
    if (n < 2) {
        return n;
    } else {
        #pragma omp task shared(left)
        left = fib(n - 1);

        #pragma omp task shared(right)
        right = fib(n - 2);

        #pragma omp taskwait
        return left + right;
    }
}

// Main
#pragma omp parallel
{
    #pragma omp single
    result = fib(n)
}

```

#### Some of the Pragma functions:

```

int omp_get_num_threads(); // to get number of threads (N)
int omp_get_thread_num(); // to get ID of the thread (0 to N - 1)
void omp_set_num_threads(int N); // to set number of threads to run

```

#### Plotting graphs using GNUPlot:

- Save the metrics using file handling operations to a textfile/csv file.
- Time the code using #include <ctime> header.

```

std::ofstream fp("metrics.txt");

double start_time, end_time;
start_time = omp_get_wtime();
#pragma omp parallel for
{
    // stmts
}
end_time = omp_get_wtime();
fp << n_value << " " << (end_time - start_time) << "\n";
fp.close();

```

#### GNUPlot code:

<Example: Plotting 2 curves (size vs serial, parallel time) using GNUPlot>

```

set terminal png size 800,600
set output 'performance_plot.png'
set title "Performance: Size vs Time"
set xlabel "Size (N)"
set ylabel "Time Taken (s)"
set gridplot 'metrics.txt' using 1:2 with linespoints title 'Serial Execution Time' \
    using 1:3 with linespoints title 'Parallel Execution Time'

```

To get graph, type `gnuplot -p plot.gp` in terminal

#### Sample code (Dot product parallelization):

```

#include <omp.h>
#include <iostream>
#include <vector>

int main() {
    const int N = 100000; // Size of the vectors
    std::vector<double> a(N, 1.0); // Vector a initialized with 1.0
    std::vector<double> b(N, 2.0); // Vector b initialized with 2.0

    double dot_product = 0.0;

    #pragma omp parallel for reduction(+:dot_product)
    for (int i = 0; i < N; i++) {
        dot_product += a[i] * b[i];
    }

    std::cout << "Dot product: " << dot_product << std::endl;
    return 0;
}

```